

Decision Trees and Random Forest Project

June 10, 2018

1 Random Forest Project

For this project we will be exploring publicly available data from [LendingClub.com](https://lendingclub.com). Lending Club connects people who need money (borrowers) with people who have money (investors). Hopefully, as an investor you would want to invest in people who showed a profile of having a high probability of paying you back. We will try to create a model that will help predict this.

Lending club had a [very interesting year in 2016](#), so let's check out some of their data and keep the context in mind. This data is from before they even went public.

We will use lending data from 2007-2010 and be trying to classify and predict whether or not the borrower paid back their loan in full. You can download the data from [here](#) or just use the csv already provided. It's recommended you use the csv provided as it has been cleaned of NA values.

Here are what the columns represent: * credit.policy: 1 if the customer meets the credit underwriting criteria of LendingClub.com, and 0 otherwise. * purpose: The purpose of the loan (takes values "credit_card", "debt_consolidation", "educational", "major_purchase", "small_business", and "all_other"). * int.rate: The interest rate of the loan, as a proportion (a rate of 11% would be stored as 0.11). Borrowers judged by LendingClub.com to be more risky are assigned higher interest rates. * installment: The monthly installments owed by the borrower if the loan is funded. * log.annual.inc: The natural log of the self-reported annual income of the borrower. * dti: The debt-to-income ratio of the borrower (amount of debt divided by annual income). * fico: The FICO credit score of the borrower. * days.with.cr.line: The number of days the borrower has had a credit line. * revol.bal: The borrower's revolving balance (amount unpaid at the end of the credit card billing cycle). * revol.util: The borrower's revolving line utilization rate (the amount of the credit line used relative to total credit available). * inq.last.6mths: The borrower's number of inquiries by creditors in the last 6 months. * delinq.2yrs: The number of times the borrower had been 30+ days past due on a payment in the past 2 years. * pub.rec: The borrower's number of derogatory public records (bankruptcy filings, tax liens, or judgments).

2 Import Libraries

```
In [20]: import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
```

2.1 Get the Data

**** Use pandas to read loan_data.csv as a dataframe called loans.****

```
In [6]: loans = pd.read_csv('loan_data.csv')
```

**** Check out the info(), head(), and describe() methods on loans.****

```
In [7]: loans.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9578 entries, 0 to 9577
Data columns (total 14 columns):
credit.policy      9578 non-null int64
purpose           9578 non-null object
int.rate          9578 non-null float64
installment       9578 non-null float64
log.annual.inc    9578 non-null float64
dti               9578 non-null float64
fico              9578 non-null int64
days.with.cr.line 9578 non-null float64
revol.bal         9578 non-null int64
revol.util        9578 non-null float64
inq.last.6mths    9578 non-null int64
delinq.2yrs       9578 non-null int64
pub.rec           9578 non-null int64
not.fully.paid    9578 non-null int64
dtypes: float64(6), int64(7), object(1)
memory usage: 1.0+ MB
```

```
In [8]: loans.head()
```

```
Out[8]:
```

	credit.policy		purpose	int.rate	installment	log.annual.inc	\
0	1	debt_consolidation	0.1189	829.10	11.350407		
1	1	credit_card	0.1071	228.22	11.082143		
2	1	debt_consolidation	0.1357	366.86	10.373491		
3	1	debt_consolidation	0.1008	162.34	11.350407		
4	1	credit_card	0.1426	102.92	11.299732		

	dti	fico	days.with.cr.line	revol.bal	revol.util	inq.last.6mths	\
0	19.48	737	5639.958333	28854	52.1	0	
1	14.29	707	2760.000000	33623	76.7	0	
2	11.63	682	4710.000000	3511	25.6	1	
3	8.10	712	2699.958333	33667	73.2	1	
4	14.97	667	4066.000000	4740	39.5	0	

	delinq.2yrs	pub.rec	not.fully.paid
0	0	0	0

1	0	0	0
2	0	0	0
3	0	0	0
4	1	0	0

In [9]: `loans.describe()`

```
Out [9]:
```

	credit.policy	int.rate	installment	log.annual.inc	dti \
count	9578.000000	9578.000000	9578.000000	9578.000000	9578.000000
mean	0.804970	0.122640	319.089413	10.932117	12.606679
std	0.396245	0.026847	207.071301	0.614813	6.883970
min	0.000000	0.060000	15.670000	7.547502	0.000000
25%	1.000000	0.103900	163.770000	10.558414	7.212500
50%	1.000000	0.122100	268.950000	10.928884	12.665000
75%	1.000000	0.140700	432.762500	11.291293	17.950000
max	1.000000	0.216400	940.140000	14.528354	29.960000

	fico	days.with.cr.line	revol.bal	revol.util \
count	9578.000000	9578.000000	9.578000e+03	9578.000000
mean	710.846314	4560.767197	1.691396e+04	46.799236
std	37.970537	2496.930377	3.375619e+04	29.014417
min	612.000000	178.958333	0.000000e+00	0.000000
25%	682.000000	2820.000000	3.187000e+03	22.600000
50%	707.000000	4139.958333	8.596000e+03	46.300000
75%	737.000000	5730.000000	1.824950e+04	70.900000
max	827.000000	17639.958330	1.207359e+06	119.000000

	inq.last.6mths	delinq.2yrs	pub.rec	not.fully.paid
count	9578.000000	9578.000000	9578.000000	9578.000000
mean	1.577469	0.163708	0.062122	0.160054
std	2.200245	0.546215	0.262126	0.366676
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	1.000000	0.000000	0.000000	0.000000
75%	2.000000	0.000000	0.000000	0.000000
max	33.000000	13.000000	5.000000	1.000000

3 Exploratory Data Analysis

**** Create a histogram of two FICO distributions on top of each other, one for each credit.policy outcome.****

In [17]: `plt.figure(figsize=(10,6))`

```
loans[loans['credit.policy']==1]['fico'].hist(bins=35,color='blue',
                                              label='Credit Policy = 1',alpha=0.6)


loans[loans['credit.policy']==0]['fico'].hist(bins=35,color='red',
```

```

plt.legend(
label='Credit Policy = 0',alpha=0.6)

```

Out[17]: <matplotlib.legend.Legend at 0x10cf4c588>



**** Create a similar figure, except this time select by the not.fully.paid column.****

```

In [18]: plt.figure(figsize=(10,6))


         loans[loans['not.fully.paid']==1]['fico'].hist(bins=35,color='blue',
                                                         label='not.fully.paid = 1',alpha=0.6)

         loans[loans['not.fully.paid']==0]['fico'].hist(bins=35,color='red',
                                                         label='not.fully.paid = 0',alpha=0.6)

         plt.legend()

```


Out[18]: <matplotlib.legend.Legend at 0x11448ca90>



**** Create a countplot using seaborn showing the counts of loans by purpose, with the color hue defined by not.fully.paid. ****

In [22]: `plt.figure(figsize=(11,7))`
`sns.countplot(x='purpose',hue='not.fully.paid',data=loans,palette='Set1')`


Out[22]: `<matplotlib.axes._subplots.AxesSubplot at 0x11470a358>`



**** Let's see the trend between FICO score and interest rate. Recreate the following jointplot.****

```
In [23]: sns.jointplot(x='fico',y='int.rate',data=loans,color='purple')
```

```
Out[23]: <seaborn.axisgrid.JointGrid at 0x1a1684f3c8>
```




**** Create the following lmplots to see if the trend differed between not.fully.paid and credit.policy.****

```
In [26]: plt.figure(figsize=(11,7))
         sns.lmplot(y='int.rate',x='fico',data=loans,hue='credit.policy',
                   col='not.fully.paid',palette='Set1')
```

```
Out[26]: <seaborn.axisgrid.FacetGrid at 0x1a17f1a240>
```

<matplotlib.figure.Figure at 0x1a17f2eeb8>



4 Setting up the Data

Let's get ready to set up our data for our Random Forest Classification Model!

```
In [27]: loans.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9578 entries, 0 to 9577
Data columns (total 14 columns):
credit.policy      9578 non-null int64
purpose           9578 non-null object
int.rate          9578 non-null float64
installment       9578 non-null float64
log.annual.inc    9578 non-null float64
dti               9578 non-null float64
fico              9578 non-null int64
days.with.cr.line 9578 non-null float64
revol.bal         9578 non-null int64
revol.util        9578 non-null float64
inq.last.6mths    9578 non-null int64
delinq.2yrs       9578 non-null int64
pub.rec           9578 non-null int64
not.fully.paid    9578 non-null int64
dtypes: float64(6), int64(7), object(1)
memory usage: 1.0+ MB
```

4.1 Categorical Features

Notice that the **purpose** column as categorical

```
In [28]: cat_feats = ['purpose']
```

```
In [29]: final_data = pd.get_dummies(loans, columns=cat_feats, drop_first=True)
```

```
In [30]: final_data.head()
```

```
Out [30]:
```

	credit.policy	int.rate	installment	log.annual.inc	dti	fico	\
0	1	0.1189	829.10	11.350407	19.48	737	
1	1	0.1071	228.22	11.082143	14.29	707	
2	1	0.1357	366.86	10.373491	11.63	682	
3	1	0.1008	162.34	11.350407	8.10	712	
4	1	0.1426	102.92	11.299732	14.97	667	

	days.with.cr.line	revol.bal	revol.util	inq.last.6mths	delinq.2yrs	\
0	5639.958333	28854	52.1	0	0	
1	2760.000000	33623	76.7	0	0	
2	4710.000000	3511	25.6	1	0	
3	2699.958333	33667	73.2	1	0	
4	4066.000000	4740	39.5	0	1	

	pub.rec	not.fully.paid	purpose_credit_card	purpose_debt_consolidation	\
0	0	0	0		1
1	0	0	1		0
2	0	0	0		1
3	0	0	0		1
4	0	0	1		0

	purpose_educational	purpose_home_improvement	purpose_major_purchase	\
0	0	0		0
1	0	0		0
2	0	0		0
3	0	0		0
4	0	0		0

	purpose_small_business
0	0
1	0
2	0
3	0
4	0

4.2 Train Test Split

**** Use sklearn to split data into a training set and a testing set.****

```
In [31]: from sklearn.cross_validation import train_test_split
```



```
/Users/Momin/anaconda3/lib/python3.6/site-packages/sklearn/cross_validation.py:41: DeprecationWarning:
  "This module will be removed in 0.20.", DeprecationWarning)
```

```
In [32]: X = final_data.drop('not.fully.paid',axis=1)
        y = final_data['not.fully.paid']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=
```

4.3 Training a Decision Tree Model

Let's start by training a single decision tree first!

**** Import DecisionTreeClassifier****

```
In [33]: from sklearn.tree import DecisionTreeClassifier
```

Create an instance of DecisionTreeClassifier() called dtree and fit it to the training data.

```
In [34]: dtree = DecisionTreeClassifier()
```

```
In [35]: dtree.fit(X_train,y_train)
```

```
Out[35]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=None,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort=False, random_state=None,
                                splitter='best')
```

4.4 Predictions and Evaluation of Decision Tree

Create predictions from the test set and create a classification report and a confusion matrix.

```
In [38]: predictions = dtree.predict(X_test)
```

```
In [39]: from sklearn.metrics import classification_report, confusion_matrix
```

```
In [40]: print(classification_report(y_test,predictions))
```

	precision	recall	f1-score	support
0	0.85	0.84	0.84	2408
1	0.20	0.21	0.21	466
avg / total	0.74	0.74	0.74	2874

```
In [41]: print(confusion_matrix(y_test,predictions))
```

```
[[2015 393]
 [ 366 100]]
```

4.5 Training the Random Forest model

Now its time to train our model!

Create an instance of the RandomForestClassifier class and fit it to our training data from the previous step.

```
In [42]: from sklearn.ensemble import RandomForestClassifier
```

```
In [43]: rfc = RandomForestClassifier(n_estimators=300)
         rfc.fit(X_train,y_train)
```

```
Out[43]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=None, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=300, n_jobs=1,
                                oob_score=False, random_state=None, verbose=0,
                                warm_start=False)
```

4.6 Predictions and Evaluation

Let's predict off the y_test values and evaluate our model.

**** Predict the class of not.fully.paid for the X_test data.****

```
In [44]: predictions = rfc.predict(X_test)
```

Now create a classification report from the results. Do you get anything strange or some sort of warning?

```
In [45]: print(classification_report(y_test,predictions))
```

	precision	recall	f1-score	support
0	0.84	0.99	0.91	2408
1	0.28	0.01	0.02	466
avg / total	0.75	0.84	0.77	2874

```
In [46]: print(confusion_matrix(y_test,predictions))
```

```
[[2395  13]
 [ 461   5]]
```

What performed better the random forest or the decision tree?

```
In [36]: # Depends what metric you are trying to optimize for.
         # Notice the recall for each class for the models.
         # Neither did very well, more feature engineering is needed.
```