

MNIST Multi-Layer Perceptron

June 10, 2018

```
In [9]: import tensorflow as tf
import matplotlib.pyplot as plt


# Import MNIST data
from tensorflow.examples.tutorials.mnist import input_data
mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)
```

```
Extracting MNIST_data/train-images-idx3-ubyte.gz
Extracting MNIST_data/train-labels-idx1-ubyte.gz
Extracting MNIST_data/t10k-images-idx3-ubyte.gz
Extracting MNIST_data/t10k-labels-idx1-ubyte.gz
```

```
In [10]: %matplotlib inline
```

```
In [12]: # visualize what the data looks like
sample = mnist.train.images[2].reshape(28,28)
plt.imshow(sample, cmap='Greys')
```

```
Out[12]: <matplotlib.image.AxesImage at 0x1c2b813710>
```



```

In [13]: # Define parameters for training
         # learning rate (how quickly the cost function is adjusted)
         learning_rate = 0.001
         # training epochs (how many training cycle to go through)
         training_epochs = 15
         # size of the batches of training data
         batch_size = 100

In [14]: # 10 possible outputs from 0 - 9
         n_classes = 10
         # Num of samples in data
         n_samples = mnist.train.num_examples

In [15]: # what we expect the input to look like
         n_input = 784

In [16]: # number of neurons in first hidden layer
         n_hidden_1 = 256
         # number of neurons in second hidden layer
         n_hidden_2 = 256

In [17]: def multilayer_perceptron(x,weights,biases):
         '''
         x: Placeholder for Data Input
         weights: Dict of weights
         biases: dict of bias values
         '''

         # First Hidden Layer with RELU Activation
         #  $X * W + B$ 
         layer_1 = tf.add(tf.matmul(x,weights['h1']), biases['b1'])
         #  $RELU(X * W + B) = RELU \rightarrow f(x) = \max(0, x)$ 
         layer_1 = tf.nn.relu(layer_1)

         #Second Hidden Layer
         layer_2 = tf.add(tf.matmul(layer_1,weights['h2']),biases['b2'])
         layer_2 = tf.nn.relu(layer_2)

         #Last Output Layer
         out_layer = tf.matmul(layer_2,weights['out']) + biases['out']

         return out_layer

In [18]: # make weight dictionary

```

```

weights = {
    'h1':tf.Variable(tf.random_normal([n_input,n_hidden_1])),
    'h2':tf.Variable(tf.random_normal([n_hidden_1,n_hidden_2])),
    'out':tf.Variable(tf.random_normal([n_hidden_2,n_classes]))
}

In [19]: # make biases dictionary
biases = {
    'b1':tf.Variable(tf.random_normal([n_hidden_1])),
    'b2':tf.Variable(tf.random_normal([n_hidden_2])),
    'out':tf.Variable(tf.random_normal([n_classes]))
}

In [20]: x = tf.placeholder('float',[None,n_input])

In [21]: y = tf.placeholder('float',[None,n_classes])

In [22]: # Construct model
pred = multilayer_perceptron(x, weights, biases)

In [25]: # Define loss and optimizer
cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=pred, labels=y))
optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(cost)

In [26]: # Run session

sess = tf.InteractiveSession()

In [28]: # Initializing the variables
init = tf.global_variables_initializer()

In [31]: # 15 loops
# training_epochs = 15

sess.run(init)

for epoch in range(training_epochs):

    # Cost
    avg_cost = 0.0

    total_batch = int(n_samples/batch_size)

    for i in range(total_batch):

        batch_x,batch_y = mnist.train.next_batch(batch_size)

        _,c = sess.run([optimizer,cost],feed_dict={x:batch_x,y:batch_y})

```

```

        # Average Loss
        avg_cost += c/total_batch

    print("Epoch: {} cost={:.4f}".format(epoch+1,avg_cost))

    print("Model has completed {} Epochs of training".format(training_epochs))

```

```

Epoch: 1 cost=156.2893
Epoch: 2 cost=38.3124
Epoch: 3 cost=24.2077
Epoch: 4 cost=16.6557
Epoch: 5 cost=12.2072
Epoch: 6 cost=9.0535
Epoch: 7 cost=6.7135
Epoch: 8 cost=5.0896
Epoch: 9 cost=3.7315
Epoch: 10 cost=2.8349
Epoch: 11 cost=2.2116
Epoch: 12 cost=1.6116
Epoch: 13 cost=1.1371
Epoch: 14 cost=0.9615
Epoch: 15 cost=0.8236
Model has completed 15 Epochs of training

```

```

In [32]: # Evaluating the model

```

```

    correct_predictions = tf.equal(tf.argmax(pred,1),tf.argmax(y,1))

```

```

In [38]: # change from bool to float

```

```

    correct_predictions = tf.cast(correct_predictions,'float')

```

```

In [39]: print(correct_predictions[1])

```

```

Tensor("strided_slice_3:0", shape=(), dtype=float32)

```

```

In [40]: accuracy = tf.reduce_mean(correct_predictions)

```

```

In [44]: accuracy.eval({x:mnist.test.images,y:mnist.test.labels})

```

```

Out[44]: 0.94919997

```